

```
In [24]: import pandas as pd
import numpy as np
from sklearn.ensemble import IsolationForest
from sklearn.preprocessing import StandardScaler
import pickle
from my_functions import (load_data, preprocess_data,
    plotting_confusion_matrix, get_features,
    calculating_marketing_campaign,
)
from sklearn.model_selection import train_test_split, StratifiedKFold, cr
from sklearn.metrics import classification_report, confusion_matrix, roc_
from catboost import CatBoostClassifier
import matplotlib.pyplot as plt
```

```
In [26]: # Load and preprocess your data as before
train_data = preprocess_data(load_data('data/churn_train_data.pkl'))
_, features, _ = load_data('CatBoostClassifier_26062025_07_17.pickle')
train_data.shape
```

Out[26]: (150000, 768)

```
In [27]: y = train_data['target']
X = train_data[features]
print('target shape', y.shape)
print('train data shape', X.shape)

# Split data FIRST
X_tr, X_te, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42,
)

# Fit scaler ONLY on X_train
scaler = StandardScaler()
scaler.fit(X_tr)
X_train_scaled = pd.DataFrame(scaler.transform(X_tr), index=X_tr.index, c
X_test_scaled = pd.DataFrame(scaler.transform(X_te), index=X_te.index, co

# Fit Isolation Forest ONLY on X_train_scaled
iso = IsolationForest(n_estimators=200, contamination=0.3, random_state=4
iso.fit(X_train_scaled)

# Compute iso_score for both splits (flip sign so higher = more anomalous
X_train = X_tr.copy()
X_test = X_te.copy()

X_train['iso_score'] = -iso.decision_function(X_train_scaled)
X_test['iso_score'] = -iso.decision_function(X_test_scaled)

distribution = lambda s: {k: float(f'{v:.4f}')} for k, v in s.value_counts
print(f"Train target class distribution: {distribution(y_train)}")
print(f"Test target class distribution: {distribution(y_test)}")
```

```
target shape (150000,)
train data shape (150000, 315)
Train target class distribution: {0.0: 0.9362, 1.0: 0.0638}
Test target class distribution: {0.0: 0.9355, 1.0: 0.0645}
```

In [28]: *# Train CatBoost with early stopping*

```
model = CatBoostClassifier(
    auto_class_weights='Balanced',
    eval_metric='F1',
    iterations=500,
    learning_rate=0.07,
    verbose=50,
    random_state=42
)

model.fit(
    X_train, y_train,
    eval_set=(X_test, y_test),
    plot=True,
    use_best_model=True,
    early_stopping_rounds=50
)

# 8. Evaluate
y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[: , 1]
print(classification_report(y_test, y_pred))
print('ROC AUC:', roc_auc_score(y_test, y_pred_proba))
```

```
MetricVisualizer(layout=Layout(align_self='stretch', height='500px'))
0:      learn: 0.7659104      test: 0.7750095 best: 0.7750095 (0)      to
tal: 90.6ms      remaining: 45.2s
50:      learn: 0.8146812      test: 0.8116007 best: 0.8119440 (48)      to
tal: 1.53s      remaining: 13.5s
100:      learn: 0.8251886      test: 0.8188054 best: 0.8191051 (94)      to
tal: 2.82s      remaining: 11.2s
150:      learn: 0.8336052      test: 0.8211995 best: 0.8215421 (139)      to
tal: 4.17s      remaining: 9.63s
200:      learn: 0.8451732      test: 0.8245561 best: 0.8254297 (192)      to
tal: 5.69s      remaining: 8.46s
Stopped by overfitting detector (50 iterations wait)
```

```
bestTest = 0.825429737
bestIteration = 192
```

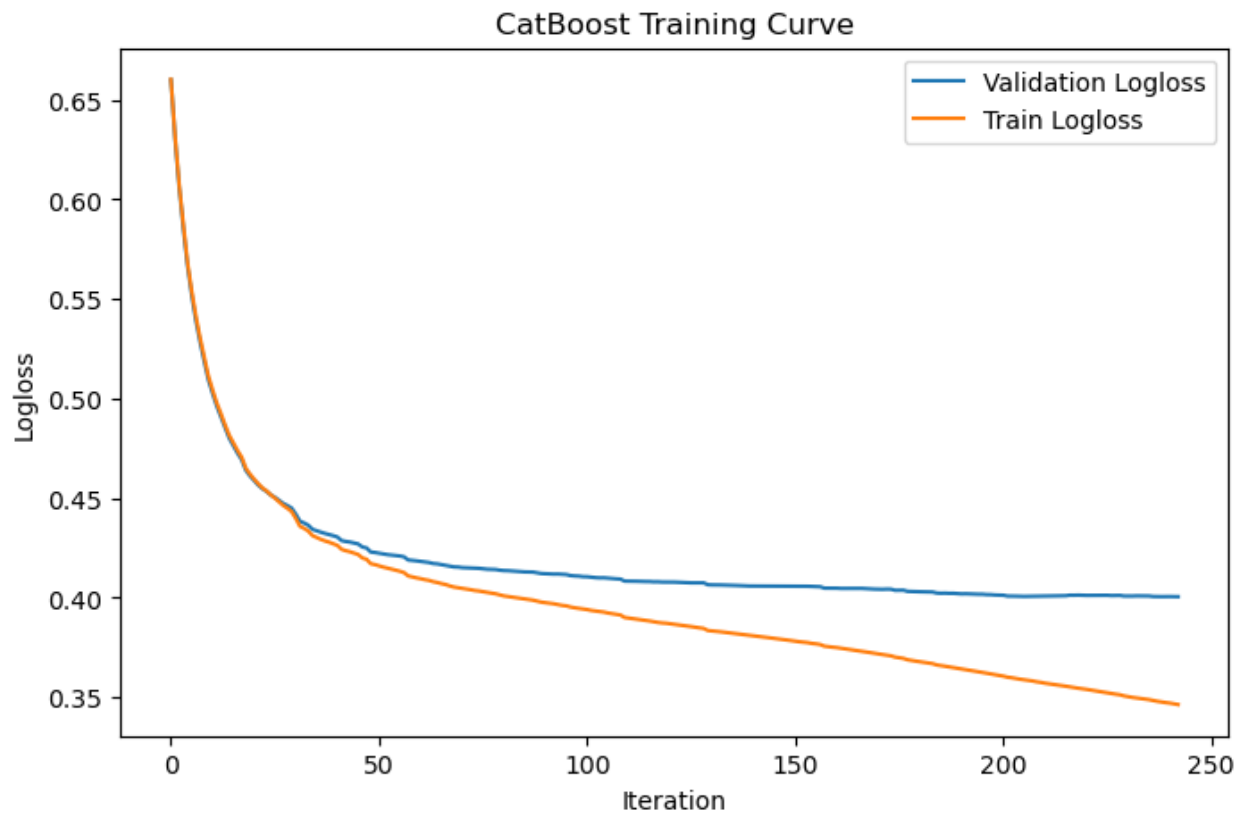
Shrink model to first 193 iterations.

	precision	recall	f1-score	support
0.0	0.98	0.85	0.91	28065
1.0	0.27	0.81	0.41	1935
accuracy			0.85	30000
macro avg	0.63	0.83	0.66	30000
weighted avg	0.94	0.85	0.88	30000

```
ROC AUC: 0.8991737619065375
```


In [30]: `evals_result = model.get_evals_result()`

```
plt.figure(figsize=(8,5))
plt.plot(evals_result['validation']['Logloss'], label='Validation Logloss')
plt.plot(evals_result['learn']['Logloss'], label='Train Logloss')
plt.xlabel('Iteration')
plt.ylabel('Logloss')
plt.title('CatBoost Training Curve')
plt.legend()
plt.show()
```



In [32]: `plotting_confusion_matrix(y_test, y_pred)`

Generating confusion matrix...



Business Impact Analysis:

True Negatives (Correctly identified retained): 23,924  
 False Positives (Incorrectly flagged as churn): 4,141  
 False Negatives (Missed churners): 377  
 True Positives (Correctly identified churn): 1,558  
 Confusion matrix analysis completed!

```
metric = F1
l_rate = 0.07
True Negatives (Correctly identified retained): 23,924
False Positives (Incorrectly flagged as churn): 4,141
False Negatives (Missed churners): 377
True Positives (Correctly identified churn): 1,558
```

```
In [34]: _, features_with_importance_gt_zero = get_features(model, X_train, y_train)


# Feature importance: CatBoost
catboost_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': model.get_feature_importance()
}).sort_values('Importance', ascending=False)

plt.figure(figsize=(8, 4))
plt.barh(features_with_importance_gt_zero['Feature'].head(10), features_w
plt.gca().invert_yaxis()
plt.title('Top 10 CatBoost Feature Importances')
plt.show()
```


Top 20 Features:

	Feature	Importance
0	loc_market_share	12.516969
1	num_act_days_mea_mnt3	4.079307
2	device_type_modem	3.878150
3	device_days_usage	3.082220
4	device_brand_sony	2.638213
5	num_act_days_mea_mnt1	2.258891
6	lt	2.201594
7	loc_is_obl_center	2.002266
8	iso_score	1.645260
9	device_type_phone	1.617112
10	Balance_uah	1.474812
11	all_cnt_std_mnt3	1.232584
12	voice_in_td_cnt_min_mnt1	1.161454
13	all_cnt_min_mnt3	1.132378
14	all_cnt_std_mnt1	1.128634
15	sms_in_cnt_max_mnt3	1.108323
16	active_ppm	1.042751
17	device_type_nan	0.982325
18	device_brand_other	0.971894
19	last_paym	0.846821

Number of Features with Importance > 0.0: 289



```
In [35]: plt.figure(figsize=(8,5))
plt.hist(X_train['iso_score'][y_train == 0], bins=50, alpha=0.6, label='N')
plt.hist(X_train['iso_score'][y_train == 1], bins=50, alpha=0.6, label='C')
plt.legend()
plt.title('iso_score Distribution by Class')
plt.show()
```



#### Analysis of the iso\_score Distribution

Both classes (Non-churn and Churn) occupy almost identical value ranges of iso\_score.

- Churners (1) do not have a distinct cluster at lower or higher values—their distribution almost fully overlaps with the majority (Non-churn, 0).
- The Churn (brown) histogram is just a tiny shadow under the much larger Non-churn curve; there's no visible region where churners dominate.

What This Means:

- iso\_score contains almost zero signal to differentiate churners from non-churners.

```
In [37]: print("Finding optimal decision threshold...")
from my_functions import calculate_optimal_threshold
# Get test predictions for threshold optimization

# Find optimal threshold on test set (better: use validation set in real
optimal_threshold, max_f1 = calculate_optimal_threshold(y_test, y_pred_proba)

# Apply optimal threshold to test probabilities to get binary predictions
y_pred_optimal = (y_pred_proba >= optimal_threshold).astype(int)

print("Threshold optimization completed!")
```

Finding optimal decision threshold...


Optimal threshold value: 0.76

Maximum F1 Score: 0.51

F1 Score vs Threshold, model CatBoost Classifier

Optimal threshold: 0.76

Max F1: 0.51

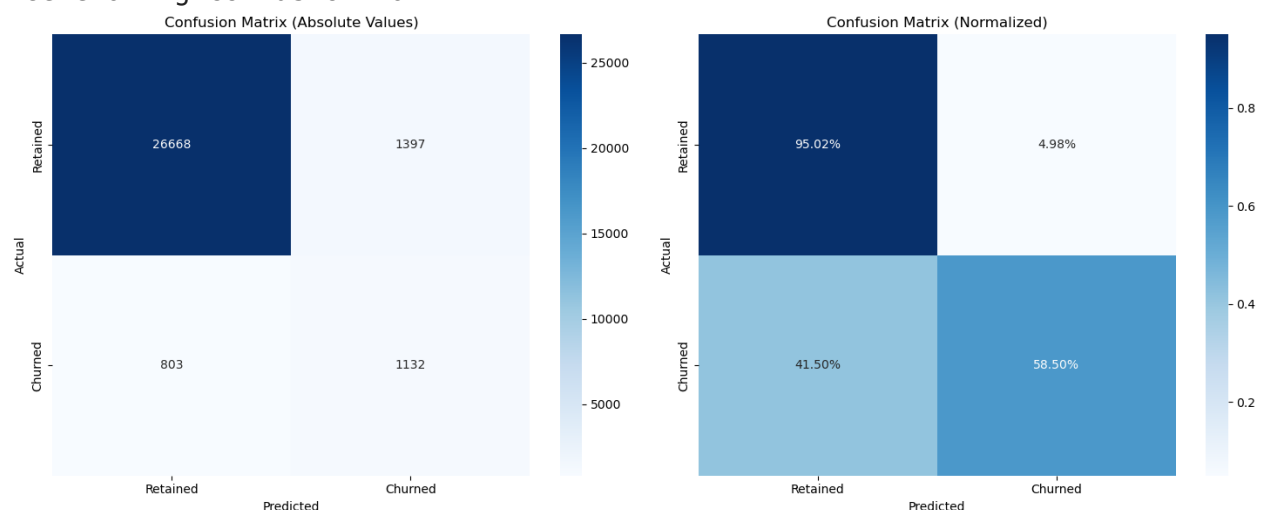


Date and time: 2025-06-26 10:17:28

Threshold optimization completed!

```
In [38]: plotting_confusion_matrix(y_test, y_pred_optimal)
```

Generating confusion matrix...



## Business Impact Analysis:

True Negatives (Correctly identified retained): 26,668  
False Positives (Incorrectly flagged as churn): 1,397  
False Negatives (Missed churners): 803  
True Positives (Correctly identified churn): 1,132  
Confusion matrix analysis completed!

TRADE-OFF between sensibility and precision

```
In [40]: # Define the stratified splitter
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Define your model
cat = CatBoostClassifier(iterations=195, learning_rate= 0.07, verbose=0,

# Use F1 as the scoring metric
scorer = make_scorer(f1_score)

# Perform cross-validation
scores = cross_val_score(cat, X, y, cv=skf, scoring=scorer)

print("F1 scores for each fold:", scores)
print("Average F1 score:", scores.mean())
```

F1 scores for each fold: [0.40014959 0.42132165 0.43529412 0.42242019 0.432531 ]

Average F1 score: 0.42234330976155404

```
In [41]: y = train_data['target']
X_scaled = pd.DataFrame(scaler.transform(X), index=X.index, columns=X.columns)
iso.fit(X_scaled)
X_final_train = X.copy()
X_final_train['iso_score'] = -iso.decision_function(X_scaled)
```

```
In [50]: # Train final optimized model
print("Training tuned model...")
final_model = CatBoostClassifier(auto_class_weights='Balanced',
                                eval_metric='AUC',
                                iterations=195,
                                verbose=50,
                                random_seed=42)

final_model.fit(X_final_train, y, early_stopping_rounds=50)
print("Final model training completed!")
```

Training tuned model...

Learning rate set to 0.391889

0:	total: 31.4ms	remaining: 6.1s
50:	total: 1.47s	remaining: 4.14s
100:	total: 2.86s	remaining: 2.66s
150:	total: 4.35s	remaining: 1.27s
194:	total: 5.71s	remaining: 0us

Final model training completed!

```
In [52]: # Load test data
test = load_data('data/churn_test_model_fe.pickle')
```

```
print("Test data structure:")
print(f"Dataset shape: {test.shape}")
```

Test data structure:  
Dataset shape: (150000, 817)

```
In [54]: # Do the same transformation as with train data
df_test = test[features].fillna(-1)
y = test['target']
X = df_test
X_scaled = pd.DataFrame(scaler.transform(df_test), index=df_test.index, c
iso.fit(X_scaled)
X_test = X.copy()
X_test['iso_score'] = -iso.decision_function(X_scaled)
```

```
In [56]: # Predict
predictions_proba = final_model.predict_proba(X_test)[:, 1]
predictions_proba
```

```
Out[56]: array([0.09169434, 0.02687977, 0.00492873, ..., 0.31815938, 0.35233475,
0.97658898])
```

```
In [57]: # Evaluate
predictions = final_model.predict(X_test)


print(classification_report(y, predictions))
print('ROC AUC:', roc_auc_score(y, predictions_proba))
```

	precision	recall	f1-score	support
0.0	0.98	0.89	0.93	140597
1.0	0.30	0.72	0.42	9403
accuracy			0.88	150000
macro avg	0.64	0.80	0.68	150000
weighted avg	0.94	0.88	0.90	150000

ROC AUC: 0.8799672640088008

```
In [58]: plotting_confusion_matrix(y, predictions)
```

Generating confusion matrix...





## Business Impact Analysis:

True Negatives (Correctly identified retained): 125,061  
False Positives (Incorrectly flagged as churn): 15,536  
False Negatives (Missed churners): 2,677  
True Positives (Correctly identified churn): 6,726  
Confusion matrix analysis completed!

In [61]: *# Calculating business impact and ROI potential for multiple retention scenarios*

```
roi_df = calculating_marketing_campaign(y, predictions, avg_customer_val
```

Calculating business impact and ROI potential for multiple retention scenarios...

### Business Impact Analysis (Retention Scenarios):

Total customers in test set: 150,000

Actual churners: 9,403.0 (6.3%)

Predicted churners: 22,262.0 (14.8%)

Correctly identified churners: 6,726

Retention Rate (%)	Potential Savings (\$)	Campaign Cost (\$)	Net ROI (\$)
ROI (%)			
-72.808373	121068.0	445240.0	-324172.0
-54.680622	201780.0	445240.0	-243460.0
-45.616746	242136.0	445240.0	-203104.0
-36.552870	282492.0	445240.0	-162748.0
-9.361243	403560.0	445240.0	-41680.0
35.958135	605340.0	445240.0	160100.0
81.277513	807120.0	445240.0	361880.0
126.596892	1008900.0	445240.0	563660.0

In [ ]:

In [ ]: