

# Vodafone Customer Churn Prediction

This project demonstrates advanced machine learning techniques to predict customer churn for Vodafone subscribers. The goal is to identify customers who are likely to leave the service, enabling proactive retention strategies with measurable business impact.

## Key Objectives:

- Build a robust churn prediction model using CatBoost
- Handle class imbalance (6.4% churn rate)
- Provide actionable business insights
- Demonstrate ROI potential for retention campaigns

## Project Structure:

1. Data Preparation and Exploration
2. Feature Engineering and Selection
3. Model Training and Optimization
4. Evaluation and Business Insights
5. Recommendations and ROI Analysis

## 1. Data Preparation and Exploration

```
In [1]: # Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pickle
import seaborn as sns
from datetime import datetime
from sklearn.metrics import (
    classification_report, confusion_matrix, f1_score,
    precision_score, recall_score, roc_auc_score, make_scorer,
)

from sklearn.model_selection import train_test_split, cross_val_score, St
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustSca
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier

import warnings
warnings.filterwarnings('ignore')

# Import custom functions
```

```

from my_functions import (
    load_data,
    save_model,
    evaluate_classification_model,
    plotting_confusion_matrix,
    drop_highly_correlated,
    get_features,
    calculate_optimal_threshold,
    plot_learning_curve,
    plot_probabilities_hist,
    plot_roc_pr_curves,
)

print("Libraries imported successfully!")
print(f"Analysis started at: {datetime.now().strftime('%Y-%m-%d %H:%M:%S')}")

```

Libraries imported successfully!  
 Analysis started at: 2025-06-23 11:30:10

```

In [2]: # Load the dataset
df_train = load_data('churn_train_data.pkl')

print("Sample data structure:")
print(f"Dataset shape: {df_train.shape}")
print(f"Memory usage: {df_train.memory_usage().sum() / 1024**2:.2f} MB")
print("Target variable: 0 = active, 1 = churned")
df_train.target

```

Sample data structure:  
 Dataset shape: (150000, 817)  
 Memory usage: 238.90 MB  
 Target variable: 0 = active, 1 = churned

```

Out[2]: 0      0.0
        1      0.0
        2      0.0
        3      0.0
        4      0.0
        ...
149995  0.0
149996  0.0
149997  0.0
149998  0.0
149999  0.0
        Name: target, Length: 150000, dtype: float16

```

```

In [3]: df_train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150000 entries, 0 to 149999
Columns: 817 entries, Ama_rchrgmnt_sum_max_mnt1 to abon_id
dtypes: float16(773), float32(1), float64(11), int8(32)
memory usage: 238.9 MB

```

```

In [4]: df_train.describe()

```

Out [4]:

|       | Ama_rchrgmnt_sum_max_mnt1 | content_clc_mea_mnt1 | content_cnt_max_mnt1 |
|-------|---------------------------|----------------------|----------------------|
| count | 150000.0                  | 150000.000000        | 150000.000000        |
| mean  | 0.0                       | 0.000000             | 0.000000             |
| std   | 0.0                       | 0.396484             | 0.000000             |
| min   | 0.0                       | 0.000000             | 0.000000             |
| 25%   | 0.0                       | 0.000000             | 9.671875             |
| 50%   | 0.0                       | 0.000000             | 11.617188            |
| 75%   | 0.0                       | 0.000000             | 13.640625            |
| max   | 0.0                       | 22.843750            | 64.062500            |

8 rows × 817 columns

In [5]: `df_train[df_train.duplicated()]`

Out [5]:

|  | Ama_rchrgmnt_sum_max_mnt1 | content_clc_mea_mnt1 | content_cnt_max_mnt1 | v |
|--|---------------------------|----------------------|----------------------|---|
|--|---------------------------|----------------------|----------------------|---|

0 rows × 817 columns

In [6]:

```
# Missing values
missing_values = df_train.isnull().sum()
missing_values_percentage = 100 * df_train.isnull().sum() / len(df_train)

missing_values_df = pd.DataFrame({'missing_values': missing_values,
                                  'percentage': missing_values_percentage})
missing_values_df.sort_values('missing_values', ascending=False).head(20)
```

Out[6]:

|                              | missing_values | percentage |
|------------------------------|----------------|------------|
| bs_of_recall_m1              | 150000         | 100.000000 |
| bs_of_succ_m1                | 150000         | 100.000000 |
| bs_of_succ_but_drop_m1       | 150000         | 100.000000 |
| bs_of_unsucc_attemp equip_m1 | 150000         | 100.000000 |
| bs_recall_rate               | 150000         | 100.000000 |
| bs_of_unsucc_low_balance_m1  | 150000         | 100.000000 |
| bs_of_attempts_all_m1        | 150000         | 100.000000 |
| bs_drop_rate                 | 150000         | 100.000000 |
| bs_succ_rate                 | 150000         | 100.000000 |
| bs_drop_call_rate            | 150000         | 100.000000 |
| device_has_gprs              | 150000         | 100.000000 |
| MV_FRAUD_BLOCK               | 149998         | 99.998667  |
| entertainment                | 149994         | 99.996000  |
| Food                         | 149989         | 99.992667  |
| tsoa_mail_cnt                | 149977         | 99.984667  |
| MV_SERV_RLH                  | 149967         | 99.978000  |
| Cars                         | 149934         | 99.956000  |
| MV_DOU_Neg_Bal               | 149887         | 99.924667  |
| Fax                          | 149845         | 99.896667  |
| Shops                        | 149475         | 99.650000  |

```
In [7]: threshold = 90.0
cols_to_drop = missing_values_df[missing_values_df['percentage'] > thresh
df_train = df_train.drop(columns=cols_to_drop)
print("Dropped columns:", cols_to_drop)
print(f"{len(cols_to_drop)} columns were dropped")
print(f"The size of our train dataset is: {df_train.shape}")
```

Dropped columns: ['DNZ\_MEAN\_days\_closed\_loan\_year2', 'DNZ\_MIN\_days\_closed\_loan\_year2', 'DNZ\_DAYS\_from\_last\_year2', 'DNZ\_MAX\_days\_closed\_loan\_year2', 'DNZ\_STD\_days\_closed\_loan\_year5', 'DNZ\_COUNT\_open\_loan\_year2', 'DNZ\_MEAN\_days\_open\_loan\_year5', 'DNZ\_COUNT\_closed\_loan\_year2', 'DNZ\_MEAN\_days\_open\_loan\_year2', 'Fax', 'tsoa\_direct\_cnt', 'tsoa\_mail\_cnt', 'SMS', 'tsoa\_chat\_cnt', 'device\_has\_gprs', 'device\_ios\_version', 'bs\_delte\_omo\_change\_tp', 'bs\_delte\_mb\_change\_tp', 'bs\_delte\_ppm\_change\_tp', 'bs\_delte\_ppd\_change\_tp', 'bs\_direct\_change\_tp', 'bs\_arpu\_change\_tp', 'bs\_day\_of\_change\_tp', 'bs\_count\_change\_tp', 'entertainment', 'Food', 'Shops', 'Cars', 'Good\_deed', 'Minutes', 'AMA', 'day\_end\_gba', 'active\_gba', 'bs\_of\_succ\_m1', 'bs\_drop\_call\_rate', 'bs\_succ\_rate', 'bs\_drop\_rate', 'bs\_of\_recall\_m1', 'bs\_of\_attempts\_all\_m1', 'bs\_of\_unsucc\_low\_balance\_m1', 'bs\_recall\_rate', 'bs\_of\_unsucc\_att emp equip\_m1', 'bs\_of\_succ\_but\_drop\_m1', 'MV\_VLR\_Guest', 'MV\_FRAUD\_BLOCK', 'MV\_SERV\_Y\_WO\_AF', 'MV\_SERV\_RLH', 'MV\_DOU\_Neg\_Bal']

48 columns were dropped

The size of our train dataset is: (150000, 769)

```
In [8]: df_train.isna().sum()
```

```
Out[8]: Ama_rchrgmnt_sum_max_mnt1      0
content_clc_mea_mnt1      0
content_cnt_max_mnt1      0
voice_out_short_part_max_mnt1      0
voice_mts_in_nrest_part_std_mnt1      0

...
MV_Migr_To      91
MV_DOU_PPM_VF      72366
MV_ot_total      41450
target      0
abon_id      0
Length: 769, dtype: int64
```

```
In [9]: df_train = df_train.fillna(-1)
df_train.isna().sum().any()
```

```
Out[9]: False
```

```
In [10]: df_train.head()
```

```
Out[10]:
```

|   | Ama_rchrgmnt_sum_max_mnt1 | content_clc_mea_mnt1 | content_cnt_max_mnt1 |
|---|---------------------------|----------------------|----------------------|
| 0 | 0                         | 0.0                  | 13.843750            |
| 1 | 0                         | 0.0                  | 11.359375            |
| 2 | 0                         | 0.0                  | 10.265625            |
| 3 | 0                         | 0.0                  | 9.976562             |
| 4 | 0                         | 0.0                  | 6.750000             |

5 rows x 769 columns

## Vizualize target distribution

```
In [11]: target_col = 'target'
```

```

# Calculate value counts and ratios
counts = df_train[target_col].value_counts()
ratios = df_train[target_col].value_counts(normalize=True)

print("Target Variable Distribution:")

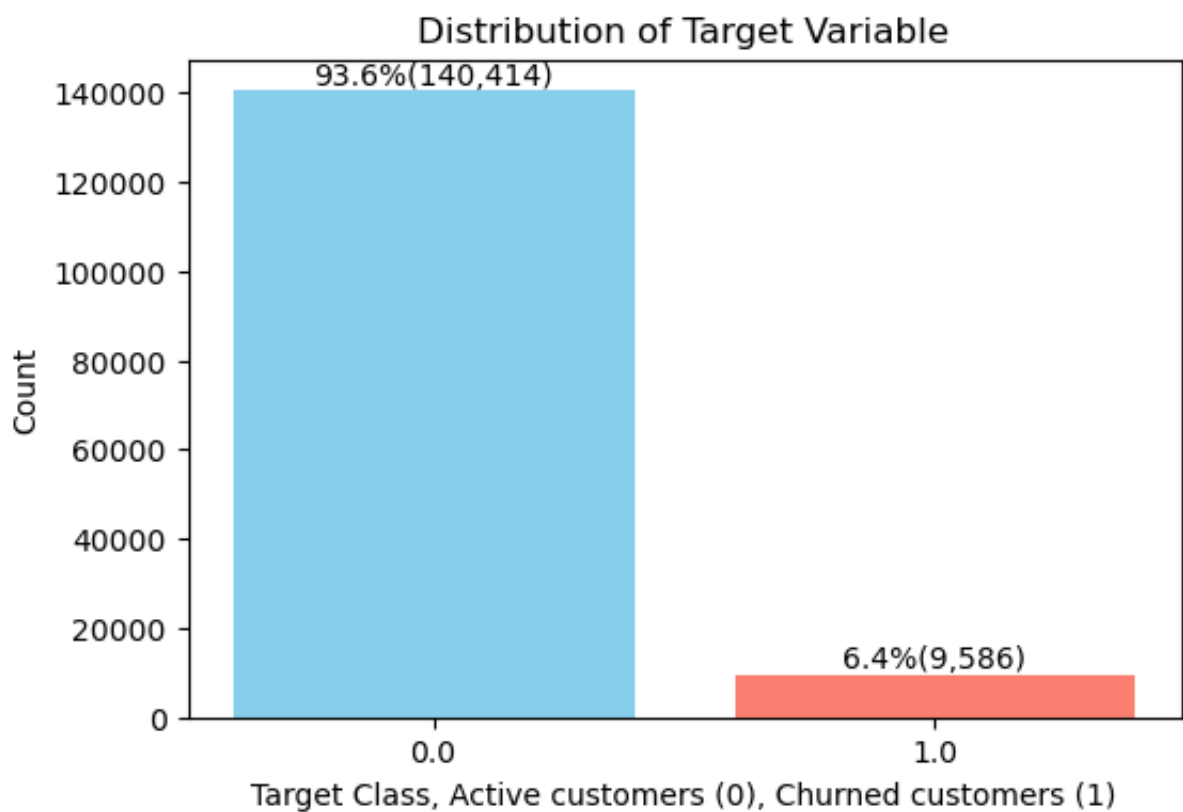
plt.figure(figsize=(6,4))
bars = plt.bar(counts.index.astype(str), counts.values, color=['skyblue',
plt.title('Distribution of Target Variable')
plt.xlabel('Target Class, Active customers (0), Churned customers (1)')
plt.ylabel('Count')

# bars with percentage and counts
for idx, value in enumerate(counts.index):
    plt.text(
        idx,
        counts[value] + max(counts.values)*0.01,
        f"{ratios[value]*100:.1f}%({counts[value]:,})",
        ha='center',
        fontsize=10
    )

plt.show()

```

Target Variable Distribution:



```

In [12]: col = 'num_act_days_min_mnt3'
plt.figure(figsize=(8, 4))
sns.histplot(df_train[col], kde=True, bins=30)
plt.title('Distribution of Number of Active Usage Days (last 3 months)')
plt.xlabel('Number of Active Days')
plt.ylabel('Frequency')

```

```

# Calculate value ratios for labeling
value_counts = df_train[col].value_counts(normalize=True)
for value, ratio in value_counts.items():
    if ratio > 0.05: # Only label values with more than 5% frequency
        plt.text(value, df_train[col].value_counts()[value], f'{ratio*100

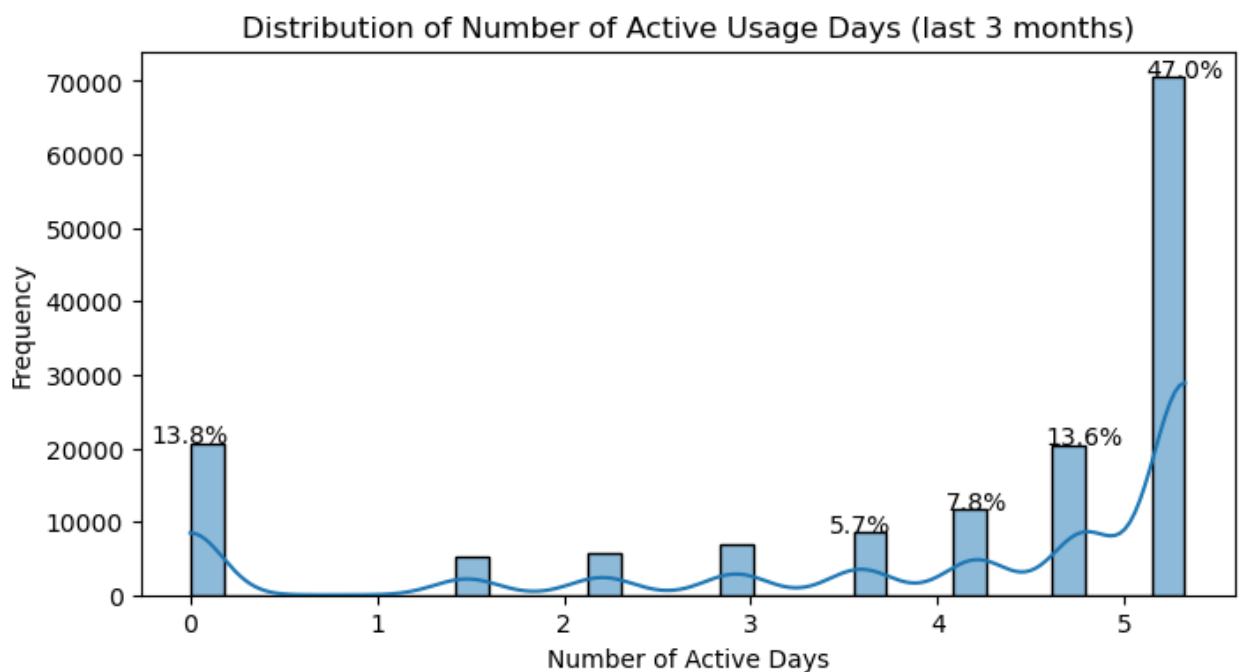
plt.show()

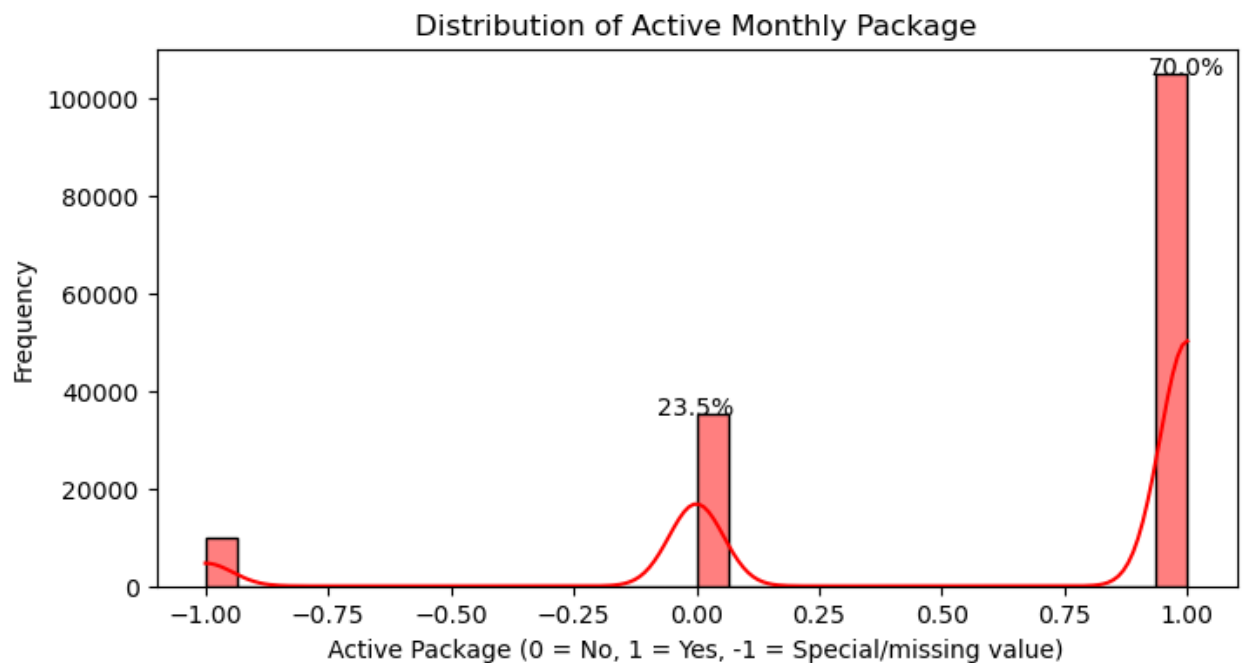
# For 'active_ppm'
col = 'active_ppm'
plt.figure(figsize=(8, 4))
sns.histplot(df_train[col], kde=True, bins=30, color='red')
plt.title('Distribution of Active Monthly Package')
plt.xlabel('Active Package (0 = No, 1 = Yes, -1 = Special/missing value)')
plt.ylabel('Frequency')

# Calculate and add ratios for 0 and 1
counts = df_train[col].value_counts(normalize=True)
for value in [0, 1]:
    if value in counts:
        plt.text(value, df_train[col].value_counts()[value], f'{counts[value]

plt.show()

```





## 2. Feature engineering and primary selection

```
In [13]: print("Removing highly correlated features...")
df_uncorr, dropped_cols = drop_highly_correlated(df_train, threshold=0.98

print(f"\nFeatures before correlation removal: {df_train.shape[1]}")
print(f"Features after correlation removal: {df_uncorr.shape[1]}")
print(f"Dropped {len(dropped_cols)} highly correlated features")
```

Removing highly correlated features...

Dropped columns: 164 nr ['pay\_avg\_max\_mnt1', 'content\_clc\_max\_mnt3', 'pay\_p2p\_in\_sum\_mea\_mnt3', 'DNZ\_COUNT\_open\_loan\_year5', 'non\_accum\_internet\_vol\_max\_mnt1', 'sms\_out\_cnt\_std\_mnt1', 'data\_3g\_tv\_cnt\_min\_mnt3', 'BS\_OVERBUNDLE\_MB\_CNT\_M1', 'all\_roam\_clc\_std\_mnt3', 'voice\_mts\_out\_nrest\_partstd\_mnt1', 'MV\_ARPU\_innet\_inc\_v\_Traf', 'voice\_out\_cmpttrs\_td\_cntstd\_mnt3', 'MV\_ap\_Roam\_d', 'MV\_ot\_innet\_out\_v', 'block\_all\_dur\_max\_mnt3', 'voice\_out\_td\_cnt\_max\_mnt1', 'voice\_out\_fix\_tar\_dur\_std\_mnt3', 'data\_3g\_dou\_mea\_mnt1', 'voice\_in\_fix\_tar\_dur\_max\_mnt3', 'ama\_volume\_max\_mnt1', 'sms\_out\_cnt\_mea\_mnt3', 'MV\_ot\_Unkn', 'conn\_in\_uniq\_cnt\_mea\_mnt1', 'MV\_ARPU\_Other\_out\_v\_Traf', 'voice\_in\_kievstar\_part\_max\_mnt3', 'voice\_in\_short\_part\_std\_mnt1', 'voice\_in\_life\_part\_max\_mnt1', 'pay\_max\_mea\_mnt1', 'sms\_clc\_max\_mnt3', 'MV\_ap\_Other', 'gprs\_tar\_vol\_mea\_mnt1', 'sms\_out\_cnt\_max\_mnt3', 'all\_clc\_std\_mnt1', 'all\_clc\_max\_mnt1', 'pay\_avg\_min\_mnt1', 'vas\_clc\_mea\_mnt3', 'all\_roam\_clc\_max\_mnt3', 'all\_home\_clc\_max\_mnt1', 'data\_3g\_tv\_cnt\_max\_mnt3', 'DNZ\_MEAN\_days\_closed\_loan\_year5', 'clc\_no\_vas\_roam\_std\_mnt1', 'block\_all\_dur\_min\_mnt3', 'gprs\_clc\_std\_mnt3', 'vas\_clc\_std\_mnt3', 'pay\_max\_std\_mnt3', 'conn\_compart\_min\_mnt3', 'voice\_in\_tar\_dur\_max\_mnt1', 'MV\_dou\_omo\_out\_v', 'block\_all\_dur\_mea\_mnt1', 'BS\_OVERBUNDLE\_MB\_CNT\_M2', 'pay\_max\_max\_mnt3', 'voice\_in\_roam\_clc\_max\_mnt1', 'voice\_out\_tar\_dur\_mea\_mnt1', 'MV\_ot\_inc\_v', 'gprs\_tar\_vol\_max\_mnt1', 'voice\_in\_cmpttrs\_td\_cnt\_std\_mnt3', 'sms\_roam\_clc\_std\_mnt3', 'gprs\_clc\_max\_mnt1', 'pay\_avg\_std\_mnt1', 'pay\_p2p\_out\_sum\_min\_mnt3', 'MV\_ot\_R\_v', 'ama\_volume\_std\_mnt3', 'all\_roam\_clc\_std\_mnt1', 'MV\_Traf\_mn\_out\_v\_Min', 'voice\_mts\_in\_nrest\_part\_std\_mnt3', 'all\_roam\_clc\_max\_mnt1', 'gprs\_tar\_vol\_mea\_mnt3', 'MV\_ot\_pstn\_out\_v', 'voice\_in\_roam\_clc\_mea\_mnt1', 'block\_all\_dur\_mea\_wk1', 'MV\_Traf\_Other\_inc\_v\_Min', 'pay\_p2p\_out\_sum\_min\_mnt1', 'data\_3g\_tv\_cnt\_min\_mnt1', 'voice\_in\_fix\_tar\_dur\_std\_mnt1', 'voice\_in\_



```
fix_tar_dur_mea_mnt1', 'voice_out_fix_tar_dur_std_mnt1', 'vas_clc_max_mnt1', 'gprs_clc_max_mnt3', 'data_3g_tar_vol_max_mnt3', 'non_accum_internet_vol_max_mnt3', 'MV_dou', 'pay_max_max_mnt1', 'pay_max_min_mnt1', 'abon_part_std_mnt1', 'all_home_clc_std_mnt3', 'pay_p2p_in_sum_std_mnt3', 'data_3g_tar_vol_std_mnt3', 'ama_volume_mea_mnt1', 'pay_p2p_out_sum_mea_mnt1', 'content_clc_std_mnt1', 'voice_out_fix_tar_dur_mea_mnt1', 'clc_no_vas_roam_max_mnt1', 'voice_in_cmpttrs_avg_durststd_mnt3', 'ama_volume_mea_mnt3', 'MV_ap_2G_d', 'voice_mts_in_nwork_part_std_mnt3', 'abon_part_max_mnt1', 'vas_clc_std_mnt1', 'pay_p2p_in_sum_max_mnt1', 'MV_ap_inc_v', 'voice_in_fix_tar_dur_std_mnt3', 'abon_part_mea_mnt3', 'voice_mts_out_nrest_partmax_mnt3', 'data_3g_tar_vol_max_mnt1', 'MV_ot_4G_d', 'MV_DOU_AP', 'MV_ARPU_Other_inc_v_Traf', 'data_3g_tv_cnt_max_mnt1', 'all_home_clc_max_mnt3', 'voice_in_roam_clc_max_mnt3', 'MV_Traf_Cont_inc_v_Min', 'pay_sum_mea_mnt3', 'accum_oth_dur_max_mnt1', 'all_home_clc_std_mnt1', 'num_act_days_min_mnt3', 'MV_ot_R_sm', 'voice_in_cmpttrs_td_cnt_mea_mnt1', 'pay_avg_mea_wk1', 'gprs_tar_vol_max_mnt3', 'pay_p2p_in_sum_mea_mnt1', 'pay_sum_max_mnt1', 'DNZ_MAX_days_closed_loan_year5', 'data_3g_tar_vol_std_mnt1', 'voice_out_short_part_std_mnt1', 'voice_mts_out_nwork_partmax_mnt1', 'MV_ap_pstn_out_v', 'voice_mts_in_nrest_part_max_mnt1', 'pay_avg_mea_mnt3', 'MV_ot_Other', 'voice_out_fix_tar_dur_mea_mnt3', 'voice_in_roam_clc_std_mnt3', 'voice_out_cmpttrs_avg_dumax_mnt3', 'abon_part_std_mnt3', 'sms_clc_max_mnt1', 'pay_sum_mea_wk1', 'voice_mts_out_nwork_partstd_mnt3', 'pay_max_std_mnt1', 'voice_out_cmpttrs_td_cntmax_mnt1', 'voice_in_life_part_max_mnt3', 'sms_out_cnt_max_mnt1', 'voice_mts_out_nrest_partmea_mnt1', 'MV_ot_Cont_v', 'voice_in_td_cnt_max_mnt1', 'accum_oth_dur_max_mnt3', 'gprs_clc_std_mnt1', 'device_type_rus_other', 'sms_roam_clc_mea_mnt3', 'MV_ot_total', 'block_all_dur_max_mnt1', 'sms_in_cnt_std_mnt3', 'DNZ_COUNT_closed_loan_year5', 'BS_OVERBUNDLE_MB_SUM_M3', 'num_act_days_min_mnt1', 'sms_clc_mea_mnt3', 'MV_ap_4G_d', 'content_clc_max_mnt1', 'content_clc_std_mnt3', 'sms_clc_mea_mnt1', 'pay_sum_mea_mnt1', 'MV_ot_Roam_d', 'MV_Traf_Roam_d_Mb', 'MV_ARPU_inc_s_Traf', 'pay_p2p_out_sum_std_mnt3', 'MV_ap_Cont_v']
```

Features before correlation removal: 769

Features after correlation removal: 605

Dropped 164 highly correlated features

```
In [14]: #df_uncorr = df_train.drop(columns=['voice_in_fix_tar_dur_mea_mnt1', 'voice_in_fix_tar_dur_std_mnt1', 'voice_in_fix_tar_dur_mea_mnt3', 'voice_in_fix_tar_dur_std_mnt3'], axis=1)
#df_uncorr.shape[1]
```

```
In [15]: # Prepare features and target
df = df_uncorr.copy()
X = df.drop(['target', 'abon_id'], axis=1, errors='ignore')
y = df['target']

print(f"Final feature matrix shape: {X.shape}")
print(f"Target shape: {y.shape}")
```

Final feature matrix shape: (150000, 603)

Target shape: (150000,)

```
In [16]: # lets use the StandardScaler
#scaler = StandardScaler()
#df_train_scaled = scaler.fit_transform(df_train)
#df_train_scaled = pd.DataFrame(df_train_scaled, index=df_train.index, columns=df_train.columns)
```

# FIRST TRAIN

```
In [17]: # Split the data into training and testing sets
print("Splitting data into training and testing sets...")
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

print(f"Training set shape: {X_train.shape}")
print(f"Testing set shape: {X_test.shape}")
print(f"Training set target distribution: {y_train.value_counts().to_dict}")
print(f"Testing set target distribution: {y_test.value_counts().to_dict()})
```

Splitting data into training and testing sets...

Training set shape: (120000, 603)

Testing set shape: (30000, 603)

Training set target distribution: {0.0: 112331, 1.0: 7669}

Testing set target distribution: {0.0: 28083, 1.0: 1917}

```
In [18]: # Count classes in your train set
neg = y_train.value_counts()[0]
pos = y_train.value_counts()[1]
scale_pos_weight = neg / pos
scale_pos_weight

# Dictionary of models to train
models = {
    "Random Forest": RandomForestClassifier(class_weight='balanced', rand
    "XGBoost": XGBClassifier(scale_pos_weight=scale_pos_weight, random_st
    "CatBoost": CatBoostClassifier(auto_class_weights='Balanced', verbose
    "Logistic Regression": LogisticRegression(class_weight='balanced', ra
    "LightGBM": LGBMClassifier(is_unbalance=True, random_state=42)
}

# Train each model and evaluate
results = []

for name, model in models.items():
    # Train the model
    model.fit(X_train, y_train)

    # Predict probabilities and binary outcomes
    prob_pred = model.predict_proba(X_test)[:, 1]
    binary_pred = model.predict(X_test)

    # Calculate metrics
    auc_score = roc_auc_score(y_test, prob_pred)
    recall = recall_score(y_test, binary_pred)
    precision = precision_score(y_test, binary_pred)
    f1 = f1_score(y_test, binary_pred)

    # Append results
    results.append({
        "Model": name,
        "AUC-ROC": auc_score,
        "F1": f1,
        "Recall": recall,
        "Precision": precision
```

```

}))

# Convert results to DataFrame
results_df = pd.DataFrame(results)
results_sorted_df = results_df.sort_values(['AUC-ROC', 'F1', 'Recall'], a
results_sorted_df

```

```

[LightGBM] [Info] Number of positive: 7669, number of negative: 112331
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.176425 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 92346
[LightGBM] [Info] Number of data points in the train set: 120000, number o
f used features: 559
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.063908 -> initscore=-2.6
84264
[LightGBM] [Info] Start training from score -2.684264

```

```

Out[18]:

```

|   | Model               | AUC-ROC  | F1       | Recall   | Precision |
|---|---------------------|----------|----------|----------|-----------|
| 4 | LightGBM            | 0.901252 | 0.415447 | 0.777256 | 0.283486  |
| 2 | CatBoost            | 0.899300 | 0.466816 | 0.706312 | 0.348610  |
| 1 | XGBoost             | 0.886481 | 0.450729 | 0.677621 | 0.337666  |
| 0 | Random Forest       | 0.878825 | 0.327532 | 0.215962 | 0.677578  |
| 3 | Logistic Regression | 0.864833 | 0.326526 | 0.775691 | 0.206786  |

## Feature importance

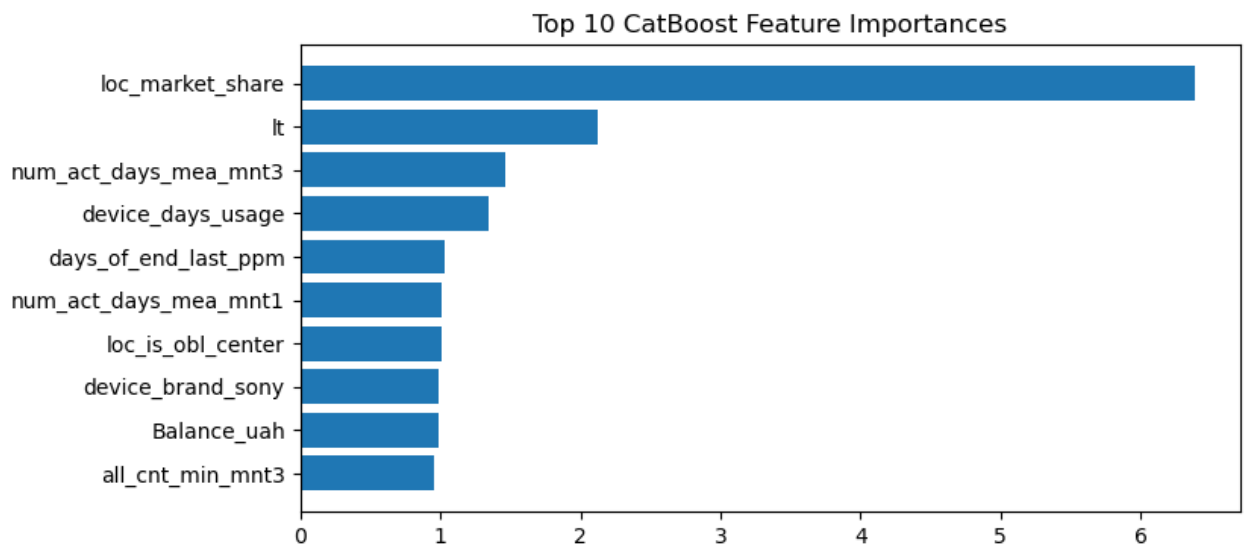
```

In [19]: # After training:
catboost_model = models["CatBoost"]

# Feature importance: CatBoost
catboost_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': catboost_model.get_feature_importance()
}).sort_values('Importance', ascending=False)

plt.figure(figsize=(8, 4))
plt.barh(catboost_importance['Feature'].head(10), catboost_importance['Im
plt.gca().invert_yaxis()
plt.title('Top 10 CatBoost Feature Importances')
plt.show()

```



```
In [20]: print("\nCatBoost Feature Importance (less important):")
print(catboost_importance[catboost_importance['Importance'] <= 0.1])
```

```
CatBoost Feature Importance (less important):
```

|     | Feature                       | Importance |
|-----|-------------------------------|------------|
| 104 | conn_out_uniq_cnt_mea_mnt3    | 0.099827   |
| 244 | voice_out_short_part_max_mnt3 | 0.099726   |
| 123 | voice_in_short_part_max_mnt3  | 0.099634   |
| 235 | clc_no_vas_roam_min_mnt1      | 0.099282   |
| 175 | conn_out_uniq_cnt_mea_mnt1    | 0.098061   |
| ..  | ...                           | ...        |
| 434 | conn_out_uniq_cnt_max_mnt1    | 0.000000   |
| 102 | Ama_rchrgmnt_sum_min_mnt3     | 0.000000   |
| 440 | Ama_rchrgmnt_sum_min_mnt1     | 0.000000   |
| 101 | content_clc_mea_mnt3          | 0.000000   |
| 0   | Ama_rchrgmnt_sum_max_mnt1     | 0.000000   |

[288 rows x 2 columns]

```
In [21]: print("Feature Importance (less important):")
col_to_drop = catboost_importance[catboost_importance['Importance'] <= 0.
df_reduced = df.drop(columns=col_to_drop)
print(f"Dropped {len(col_to_drop)} columns. New shape: {df_reduced.shape}")
```

```
Feature Importance (less important):
Dropped 288 columns. New shape: (150000, 317)
```

```
In [22]: X = df_reduced.drop(['target', 'abon_id'], axis=1, errors='ignore')
y = df_reduced['target']

# SNew split
print("Splitting data into training and testing sets...")
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

print(f"Training set shape: {X_train.shape}")
print(f"Testing set shape: {X_test.shape}")
print(f"Training set target distribution: {y_train.value_counts().to_dict}")
print(f"Testing set target distribution: {y_test.value_counts().to_dict()")
```

Splitting data into training and testing sets...  
Training set shape: (120000, 315)  
Testing set shape: (30000, 315)  
Training set target distribution: {0.0: 112331, 1.0: 7669}  
Testing set target distribution: {0.0: 28083, 1.0: 1917}

## Train again on important features

```
In [ ]: #Train again
model = CatBoostClassifier(auto_class_weights='Balanced', eval_metric='AU

model.fit(X_train, y_train, eval_set=(X_test, y_test), use_best_model=True

# Predict probabilities and binary outcomes
prob_pred = model.predict_proba(X_test)[:, 1]
binary_pred = model.predict(X_test)

results_df = evaluate_classification_model(
    y_true=y_test,
    y_pred_binary=binary_pred,
    y_pred_proba=prob_pred,
    model_name='catboost'
)
```

Learning rate set to 0.13977

|       |                                       |               |            |
|-------|---------------------------------------|---------------|------------|
| 0:    | test: 0.8413138 best: 0.8413138 (0)   | total: 31.8ms | remaining: |
| 15.9s |                                       |               |            |
| 10:   | test: 0.8817105 best: 0.8817105 (10)  | total: 335ms  | remaining: |
| 14.9s |                                       |               |            |
| 20:   | test: 0.8915978 best: 0.8915978 (20)  | total: 667ms  | remaining: |
| 15.2s |                                       |               |            |
| 30:   | test: 0.8953196 best: 0.8953196 (30)  | total: 973ms  | remaining: |
| 14.7s |                                       |               |            |
| 40:   | test: 0.8976419 best: 0.8976419 (40)  | total: 1.27s  | remaining: |
| 14.2s |                                       |               |            |
| 50:   | test: 0.8991662 best: 0.8991662 (50)  | total: 1.63s  | remaining: |
| 14.4s |                                       |               |            |
| 60:   | test: 0.9000271 best: 0.9000271 (60)  | total: 2.02s  | remaining: |
| 14.6s |                                       |               |            |
| 70:   | test: 0.9001630 best: 0.9003464 (63)  | total: 2.32s  | remaining: |
| 14s   |                                       |               |            |
| 80:   | test: 0.9006937 best: 0.9007923 (79)  | total: 2.61s  | remaining: |
| 13.5s |                                       |               |            |
| 90:   | test: 0.9013607 best: 0.9014487 (89)  | total: 2.92s  | remaining: |
| 13.1s |                                       |               |            |
| 100:  | test: 0.9014217 best: 0.9014487 (89)  | total: 3.21s  | remaining: |
| 12.7s |                                       |               |            |
| 110:  | test: 0.9012560 best: 0.9018015 (105) | total: 3.5s   | remaining: |
| 12.3s |                                       |               |            |
| 120:  | test: 0.9013372 best: 0.9018015 (105) | total: 3.8s   | remaining: |
| 11.9s |                                       |               |            |

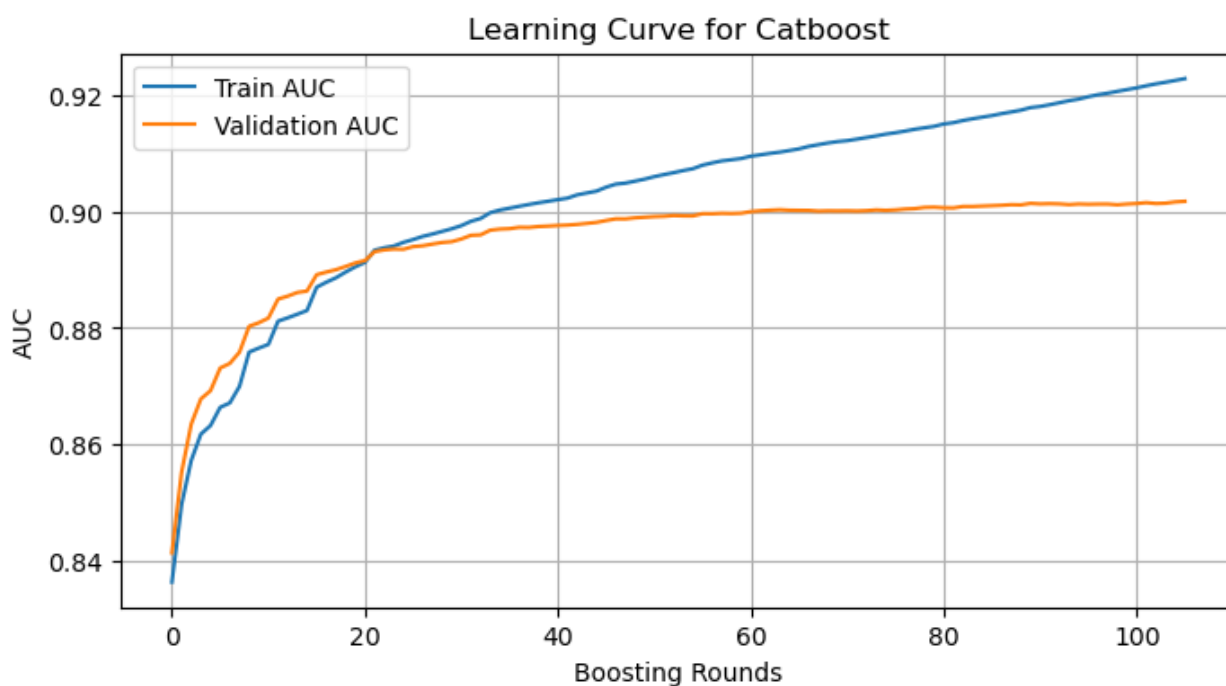
Stopped by overfitting detector (20 iterations wait)

bestTest = 0.9018015027

bestIteration = 105

Shrink model to first 106 iterations.

```
In [26]: # plot loss on train and validation
plot_learning_curve(model, X_train, y_train, X_test, y_test)
```



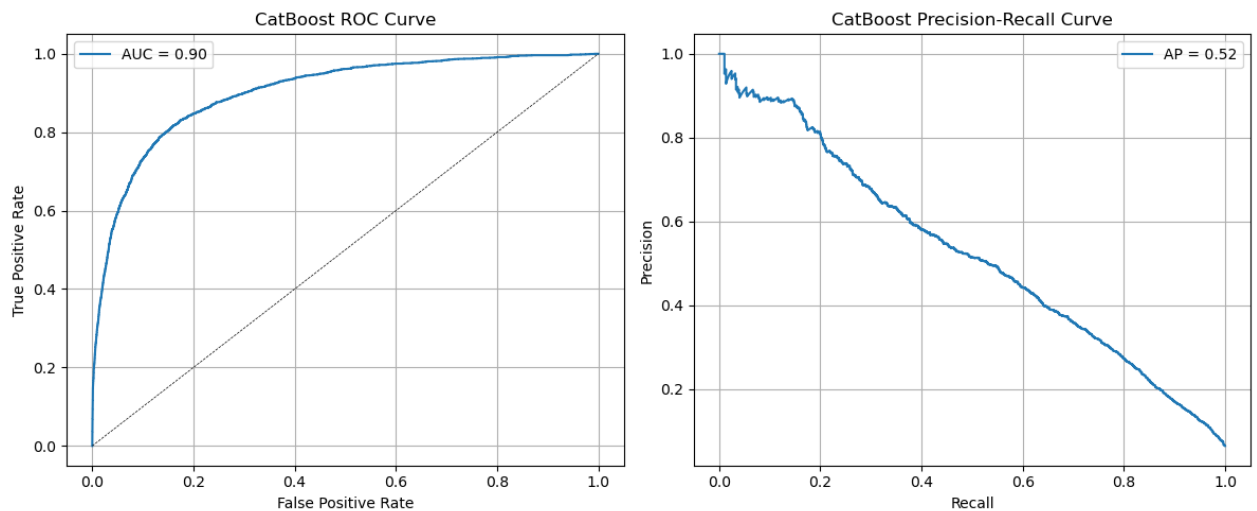
## 4. Model Evaluation

```
In [27]: # Plot ROC curves for comparison
print("Generating ROC curve comparison...")

plot_roc_pr_curves(y_test, prob_p, title_prefix='CatBoost')

print("ROC curve analysis completed!")
```

Generating ROC curve comparison...



ROC curve analysis completed!

```
In [28]: # Analyze feature importance
print("Analyzing feature importance...")

# Get top features and all important features
top_features, important_features = get_features(model, X_train, y_train,

print(f"\nTop 20 most important features identified")
print(f"Total features with importance > 0: {len(important_features)}")

# Visualize top features
feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': model.feature_importances_
}).sort_values('Importance', ascending=False).head(20)

plt.figure(figsize=(12, 8))
plt.barh(range(len(feature_importance)), feature_importance['Importance'])
plt.yticks(range(len(feature_importance)), feature_importance['Feature'])
plt.xlabel('Feature Importance')
plt.title('Top 20 Most Important Features for Churn Prediction')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

print("Feature importance analysis completed!")
```

Analyzing feature importance...

Feature Importance:

|     | Feature                    | Importance |
|-----|----------------------------|------------|
| 257 | loc_market_share           | 14.991926  |
| 200 | num_act_days_mea_mnt3      | 4.234254   |
| 265 | device_type_nan            | 3.598976   |
| 278 | device_type_phone          | 3.347024   |
| 274 | device_type_modem          | 3.016326   |
| ..  | ...                        | ...        |
| 112 | all_cnt_td_mnt3            | 0.000000   |
| 238 | data_3g_dou_std_mnt1       | 0.000000   |
| 237 | voice_out_tar_dur_std_mnt3 | 0.000000   |
| 47  | data_3g_tar_vol_mea_mnt1   | 0.000000   |
| 102 | all_cnt_mea_mnt3           | 0.000000   |

[315 rows x 2 columns]

Top 20 Features:

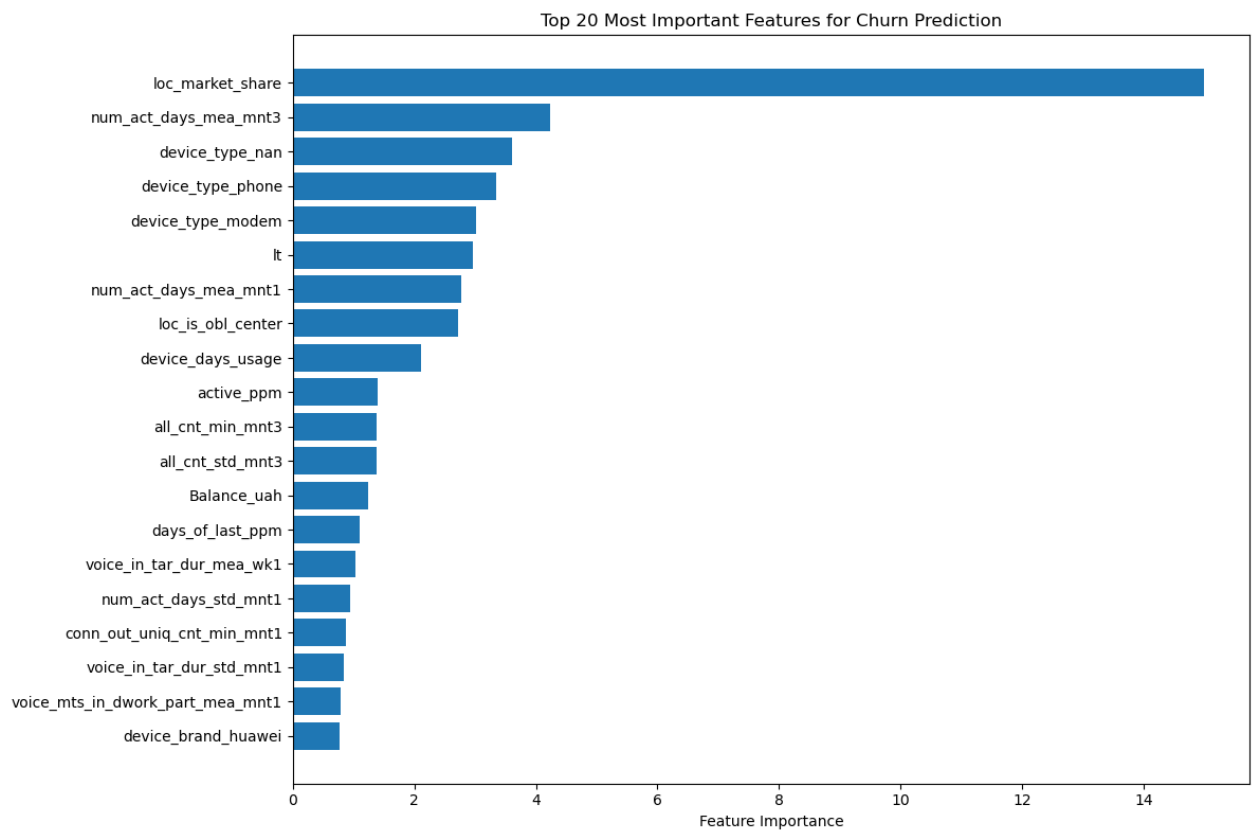
|     | Feature                          | Importance |
|-----|----------------------------------|------------|
| 257 | loc_market_share                 | 14.991926  |
| 200 | num_act_days_mea_mnt3            | 4.234254   |
| 265 | device_type_nan                  | 3.598976   |
| 278 | device_type_phone                | 3.347024   |
| 274 | device_type_modem                | 3.016326   |
| 251 | lt                               | 2.959417   |
| 203 | num_act_days_mea_mnt1            | 2.770951   |
| 255 | loc_is_obl_center                | 2.718872   |
| 259 | device_days_usage                | 2.116950   |
| 248 | active_ppm                       | 1.397445   |
| 148 | all_cnt_min_mnt3                 | 1.385086   |
| 229 | all_cnt_std_mnt3                 | 1.375491   |
| 249 | Balance_uah                      | 1.243001   |
| 247 | days_of_last_ppm                 | 1.104200   |
| 217 | voice_in_tar_dur_mea_wk1         | 1.029268   |
| 147 | num_act_days_std_mnt1            | 0.951006   |
| 115 | conn_out_uniq_cnt_min_mnt1       | 0.870035   |
| 19  | voice_in_tar_dur_std_mnt1        | 0.839326   |
| 104 | voice_mts_in_dwork_part_mea_mnt1 | 0.788617   |
| 269 | device_brand_huawei              | 0.767560   |

Number of Features with Importance > 0.0: 245

Top 20 most important features identified

Total features with importance > 0: 245





Feature importance analysis completed!

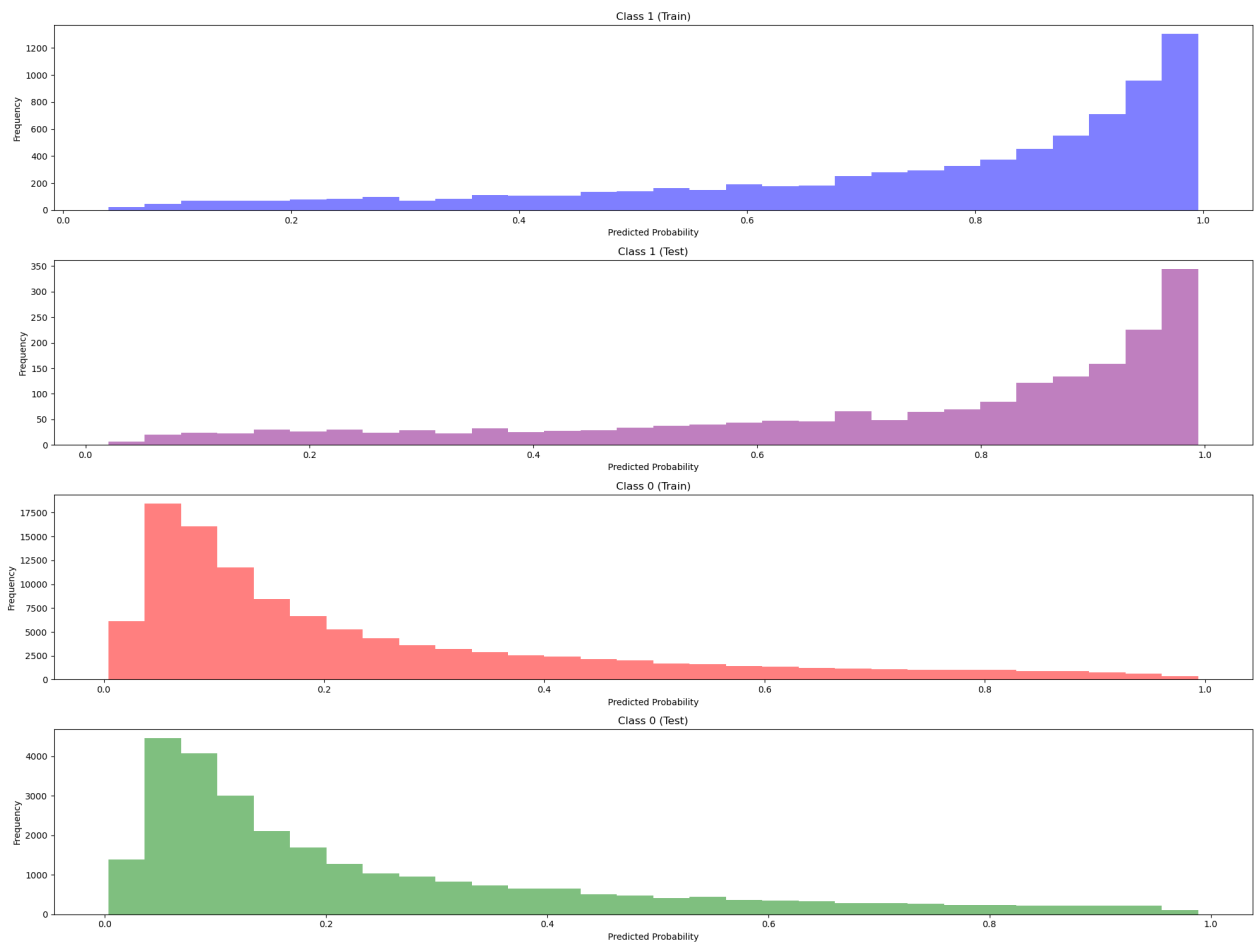
```
In [29]: # Plot probability distributions
print("Generating probability distribution plots...")

# For training data (probability of class 1)
y_train_pred_proba = model.predict_proba(X_train)[: , 1]
y_test_pred_proba = model.predict_proba(X_test)[: , 1]

# Get probability predictions for both classes
train_proba_1 = y_train_pred_proba[y_train == 1]
test_proba_1 = y_test_pred_proba[y_test == 1]
train_proba_0 = y_train_pred_proba[y_train == 0]
test_proba_0 = y_test_pred_proba[y_test == 0]

plot_probabilities_hist(train_proba_1, test_proba_1, train_proba_0, test_
print("Probability distribution analysis completed!")
```

Generating probability distribution plots...



Probability distribution analysis completed!

## Model Optimization Finding optimal decision threshold

```
In [30]: print("Finding optimal decision threshold...")

# Get test predictions for threshold optimization
catboost_test_probs = model.predict_proba(X_test)[: , 1]

# Find optimal threshold on test set (better: use validation set in real
optimal_threshold, max_f1 = calculate_optimal_threshold(y_test, catboost_

# Apply optimal threshold to test probabilities to get binary predictions
y_pred_optimal = (catboost_test_probs >= optimal_threshold).astype(int)

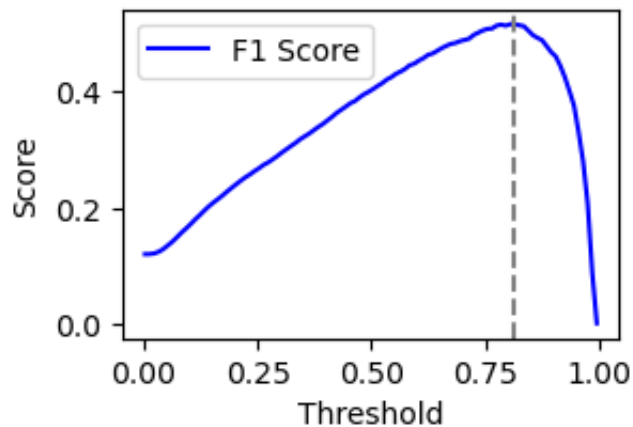
print("Threshold optimization completed!")
```

Finding optimal decision threshold...  
 Optimal threshold value: 0.81  
 Maximum F1 Score: 0.52

## F1 Score vs Threshold, model CatBoost Classifier

Optimal threshold: 0.81

Max F1: 0.52



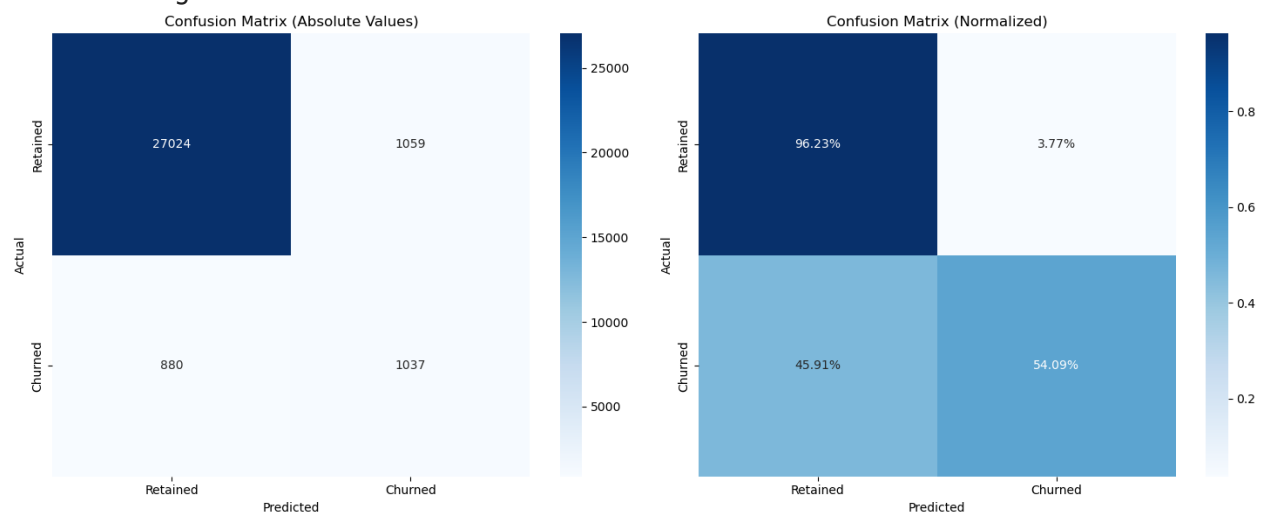
Date and time: 2025-06-23 11:40:56

Threshold optimization completed!

## Confusion matrix

```
In [31]: plotting_confusion_matrix(y_test, y_pred_optimal)
```

Generating confusion matrix...



Business Impact Analysis:

True Negatives (Correctly identified retained): 27,024

False Positives (Incorrectly flagged as churn): 1,059

False Negatives (Missed churners): 880

True Positives (Correctly identified churn): 1,037

Confusion matrix analysis completed!

```
In [32]: print('Model Evaluation')
print(classification_report(y_test, y_pred_optimal))
```

## Model Evaluation

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0          | 0.97      | 0.96   | 0.97     | 28083   |
| 1.0          | 0.49      | 0.54   | 0.52     | 1917    |
| accuracy     |           |        | 0.94     | 30000   |
| macro avg    | 0.73      | 0.75   | 0.74     | 30000   |
| weighted avg | 0.94      | 0.94   | 0.94     | 30000   |

```
In [33]: # Evaluate
results_df = evaluate_classification_model(
    y_true=y_test,
    y_pred_binary=y_pred_optimal,
    y_pred_proba=catboost_test_probs,
    model_name='catboost'
)
```

## CROSS VALIDATION

```
In [34]: # Define the stratified splitter
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Define your model
cat = CatBoostClassifier(iterations=106, verbose=0, random_seed=42)

# Use F1 as the scoring metric
scorer = make_scorer(f1_score)

# Perform cross-validation
scores = cross_val_score(cat, X, y, cv=skf, scoring=scorer)

print("F1 scores for each fold:", scores)
print("Average F1 score:", scores.mean())
```

F1 scores for each fold: [0.41918665 0.44177259 0.44528562 0.44459644 0.44657534]

Average F1 score: 0.4394833282328839

## Final train on all train data

```
In [57]: # Train final optimized model
print("Training tuned model...")
final_model = CatBoostClassifier(auto_class_weights='Balanced', eval_metric='logloss')
final_model.fit(X, y)
print("Final model training completed!")
```

```
Training tuned model...
Learning rate set to 0.5
0:      total: 32.2ms    remaining: 1.58s
10:     total: 337ms    remaining: 1.2s
20:     total: 629ms    remaining: 868ms
30:     total: 915ms    remaining: 561ms
40:     total: 1.2s     remaining: 263ms
49:     total: 1.52s    remaining: 0us
Final model training completed!
```

## Save the final model

```
In [58]: print("Saving the final model...")
save_model(final_model, list(X.columns), optimal_threshold)
print("Model saved successfully!")
print(f"Analysis completed at: {datetime.now().strftime('%Y-%m-%d %H:%M:%S')}")
```

```
Saving the final model...
Save CatBoostClassifier_23062025_12_05.pickle
Model saved successfully!
Analysis completed at: 2025-06-23 12:05:46
```

## PREDICTIONS ON TEST DATA

```
In [59]: # Load saved data
model, feature_names, threshold = load_data('CatBoostClassifier_22062025_12_05.pickle')
```

```
In [60]: # Load test data
test = load_data('churn_test_model_fe.pickle')
print("Test data structure:")
print(f"Dataset shape: {test.shape}")
```

```
Test data structure:
Dataset shape: (150000, 817)
```

```
In [61]: # Do the same transformation as with train data
df_test = test[feature_names].fillna(-1)
y = test['target']
X = df_test
# Predict
predictions_proba = model.predict_proba(X)[:, 1]
predictions_binary = (predictions_proba >= threshold).astype(int)
predictions_binary
```

```
Out[61]: array([0, 0, 0, ..., 0, 0, 1])
```

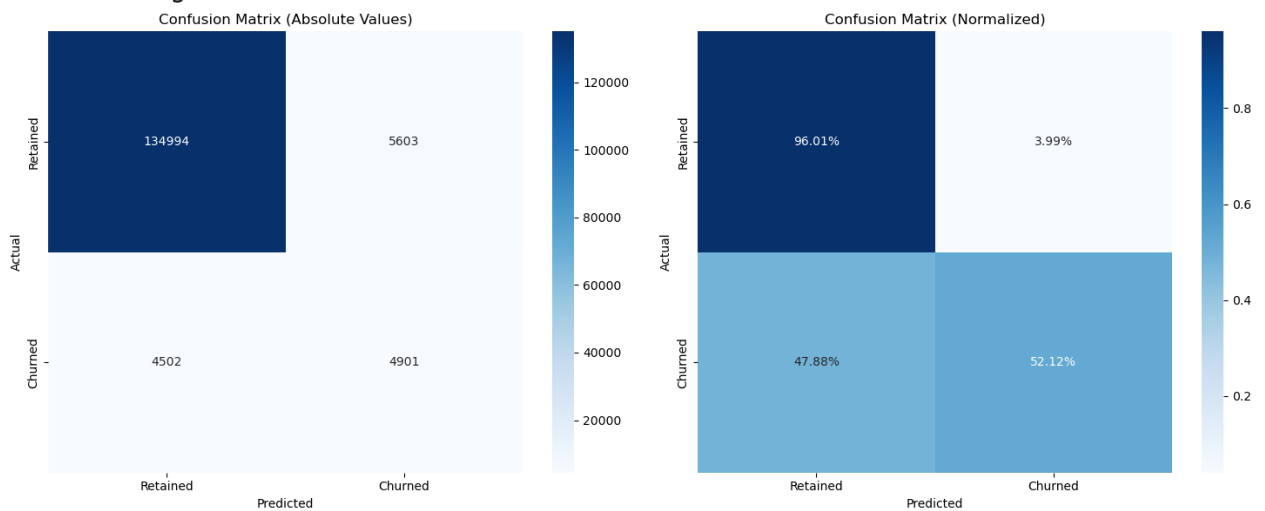
```
In [67]: # Evaluate
results_df = evaluate_classification_model(
    y_true=y,
    y_pred_binary=predictions_binary,
    y_pred_proba=predictions_proba,
    model_name='catboost'
)
results_df
```

```
Out[67]:
```

|   | Model    | AUC-ROC  | Recall   | F1      | Precision |
|---|----------|----------|----------|---------|-----------|
| 0 | catboost | 0.880256 | 0.521217 | 0.49239 | 0.466584  |

```
In [63]: plotting_confusion_matrix(y, predictions_binary)
```

Generating confusion matrix...



Business Impact Analysis:

True Negatives (Correctly identified retained): 134,994  
 False Positives (Incorrectly flagged as churn): 5,603  
 False Negatives (Missed churners): 4,502  
 True Positives (Correctly identified churn): 4,901  
 Confusion matrix analysis completed!

```
In [64]: print('Model Evaluation')
print(classification_report(y, predictions_binary))
```

Model Evaluation

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0          | 0.97      | 0.96   | 0.96     | 140597  |
| 1.0          | 0.47      | 0.52   | 0.49     | 9403    |
| accuracy     |           |        | 0.93     | 150000  |
| macro avg    | 0.72      | 0.74   | 0.73     | 150000  |
| weighted avg | 0.94      | 0.93   | 0.93     | 150000  |

## 5. Business Recommendations and ROI Analysis

```
In [65]: print("Calculating business impact and ROI potential for multiple retenti

# Metrics from the model/results
total_customers = len(y)
actual_churners = sum(y)
predicted_churners = sum(predictions_binary)
correctly_identified = sum((y == 1) & (predictions_binary == 1))

# Business assumptions
avg_customer_value = 50 # Monthly revenue per customer
```

```

campaign_cost_per_customer = 20 # Cost of retention campaign per customer

# Retention rates to analyze (from 3% to 30%)
retention_rates = [0.03, 0.05, 0.0715, 0.1, 0.15, 0.2, 0.25]

roi_results = []

for rate in retention_rates:
    potential_savings = correctly_identified * avg_customer_value * 12 *
    campaign_cost = predicted_churners * campaign_cost_per_customer
    net_roi = potential_savings - campaign_cost
    roi_percentage = (net_roi / campaign_cost) * 100 if campaign_cost > 0

    roi_results.append({
        "Retention Rate (%)": f"{rate*100:.0f}%",
        "Potential Savings ($)": potential_savings,
        "Campaign Cost ($)": campaign_cost,
        "Net ROI ($)": net_roi,
        "ROI (%)": roi_percentage
    })

roi_df = pd.DataFrame(roi_results)

print(f"\nBusiness Impact Analysis (Retention Scenarios):")
print(f"Total customers in test set: {total_customers:,}")
print(f"Actual churners: {actual_churners:,} ({actual_churners/total_cust")
print(f"Predicted churners: {predicted_churners:,} ({predicted_churners/t")
print(f"Correctly identified churners: {correctly_identified:,}\n")

print(roi_df.to_string(index=False))

```

Calculating business impact and ROI potential for multiple retention scenarios...

Business Impact Analysis (Retention Scenarios):

Total customers in test set: 150,000

Actual churners: 9,403.0 (6.3%)

Predicted churners: 10,504 (7.0%)

Correctly identified churners: 4,901

| Retention Rate (%) | Potential Savings (\$) | Campaign Cost (\$) | Net ROI (\$) | ROI (%)    |
|--------------------|------------------------|--------------------|--------------|------------|
| 3%                 | 88218.0                | 210080             | -121862.0    | -58.007426 |
| 5%                 | 147030.0               | 210080             | -63050.0     | -30.012376 |
| 7%                 | 210252.9               | 210080             | 172.9        | 0.082302   |
| 10%                | 294060.0               | 210080             | 83980.0      | 39.975248  |
| 15%                | 441090.0               | 210080             | 231010.0     | 109.962871 |
| 20%                | 588120.0               | 210080             | 378040.0     | 179.950495 |
| 25%                | 735150.0               | 210080             | 525070.0     | 249.938119 |

**Business Impact Reflection**

This analysis shows how leveraging modern data tools can directly enhance our business outcomes. By running a test campaign, we can measure actual retention rates and assess the financial impact - even when initial costs are incurred. Importantly, as retention success rates exceed 10%, the campaign shifts into profitability, delivering a strong ROI and substantial net gains for the business.

It's also important to remember that retaining an existing client is up to six times less expensive than acquiring a new one, underscoring the value of investing in effective retention strategies.

## Summary and Key Findings

### Model Performance on unseen data

- **ROC AUC: 0.88** - Excellent discrimination ability due to hard unbalanced data
- **F1 Score: 0.49** - Good balance between precision and recall
- **Optimal Threshold: 0.81** - Maximizes business value

### Churn Prevention: The Precision-Recall Trade-Off

When deploying a churn prediction model, there is always a fundamental trade-off between campaign costs and over-targeting:

- But ! the cost of losing a customer is high relative to the cost of a retention offer, it may be better to maximize recall, even if that means higher campaign costs.
- If budget constraints are critical, you may prefer to maximize precision and accept some churn.